

Uncertainty-Aware Human Querying for Knowledge-Based Robot Planning

Changmin Park¹, Seongah Park¹, Nayoung Kim¹, and Sungjin Park¹
Geonyeong Ko¹, Junho Lee¹, and Changjoo Nam¹

¹Sogang University, Seoul, Republic of Korea

cmp9877@gmail.com

Abstract

Robots operating in real environments frequently encounter uncertainty that cannot be resolved through autonomous perception or planning alone. Existing human-in-the-loop systems often address this problem through continuous supervision, failure-triggered intervention, or fixed query policies, all of which either increase operator burden or fail to request assistance before errors occur. This paper presents an integrated robot system that decides both when it is uncertain and what it should ask a human during task execution. The system maintains a symbolic knowledge base together with a particle belief over possible world states, performs online planning with POMCP, and updates the belief after each action using the received observation. When the observation leaves multiple transition outcomes plausible, the system computes entropy-based confidence, selects a task-relevant symbolic hypothesis whose answer is expected to maximally reduce belief uncertainty, queries the human through a lightweight interface, and updates the knowledge base from the response. This enables the robot to request targeted feedback after ambiguous observations rather than relying on continuous monitoring, always querying, or failure-only intervention. We evaluate the system in simulated waste sorting and tomato harvesting domains, with additional real-robot validation in tomato harvesting, and conduct a user study to measure human workload. The intended findings are that the proposed system preserves task success while using substantially fewer queries than an always-query system, and that this reduction in unnecessary interaction leads to lower subjective workload for users.

Keywords: belief-space planning, human-robot interaction, knowledge-based robotics, POMDP, proactive querying, planning under uncertainty

1 Introduction

Recent advances in robotic systems have improved autonomy in perception, manipulation, and planning, yet robots still fail in real-world environments due to perception errors, action execution failures, partial observability, and dynamic changes [22, 21, 26]. These failures often arise because the robot must select actions from an incomplete or ambiguous estimate of the current world

Figure Placeholder

fig:overview

Figure 1: Overview of the proposed framework. The final paper should show the closed loop among sensing, symbolic belief update, POMCP planning, human query, and knowledge-base commitment.

state. In tasks such as tomato harvesting and waste sorting, uncertainty about object states or categories may remain even after sensing, making human feedback a useful source of task-relevant information.

Existing ways of incorporating humans into robot autonomy remain limited. Continuous-supervision and human-on-the-loop systems rely on operators to monitor execution and detect failures [13, 10, 47, 39], which can increase workload and fatigue [4, 11, 55]. Failure-triggered systems reduce monitoring but ask only after errors occur, while fixed or uncertainty-threshold query systems can ask too often [44, 31, 38]. These approaches leave two coupled questions unresolved: *when* should the robot ask a human, and *what* should it ask?

We trigger human feedback from the robot’s belief update rather than treating it as continuous supervision or failure recovery. After each action, the robot compares the observation with possible transition outcomes encoded in its symbolic belief. If multiple task-relevant successor states remain plausible, the system estimates uncertainty using entropy and asks about the symbolic hypothesis expected to reduce the ambiguity. This follows the broader goal of selective assistance, where robots request human input only when it is useful [45, 3, 29].

This paper proposes a knowledge-based belief-space planning framework for proactive human querying. The robot maintains a symbolic knowledge base [8] with STRIPS-style state and action structure [12], while online decision making is performed using a POMDP formulation and POMCP-based planning [21, 49]. Instead of committing uncertain perceptual outputs as deterministic facts, the robot maintains a particle belief over symbolic possible worlds. After each observation, the system updates the belief over reachable frontier states, computes entropy-based confidence, and either commits the most likely state to the knowledge base or asks the human a targeted question before replanning.

The framework differs from continuous supervision, failure-triggered intervention, always-query policies, and action-level query systems: the user is not asked to monitor every decision or choose robot actions, but instead answers targeted questions about uncertain symbolic state facts.

We evaluate the framework in waste sorting and tomato harvesting, including real-robot tomato experiments. We compare against no-query, always-query, and random-query baselines, analyze sensitivity to the confidence threshold τ , and conduct a user study to examine whether selective

state-level querying reduces subjective workload.

The contributions and findings of this paper are:

- We present an integrated robot system that determines both *when* it is uncertain and *what* it should ask the human during task execution, by combining symbolic knowledge, belief-space planning, entropy-based uncertainty estimation, and hypothesis-level feedback.
- We show that the proposed system substantially reduces the number of human queries compared with an always-query system while maintaining comparable task success, demonstrating that selective querying can preserve task performance without continuous supervision.
- We show through a user study that reducing unnecessary queries also reduces human workload, supporting the central goal of making uncertainty-aware robot assistance easier for people to use.
- We validate the system across simulated and real robotic domains and compare it against no-query, always-query, and random-query baselines to isolate the effect of deciding when and what to ask.

2 Related Work

In this section, we review prior work on knowledge-based robot planning, belief-space planning under uncertainty, human involvement in robot autonomy, and uncertainty-aware query systems.

2.1 Knowledge-Based Robot Planning

Knowledge-based robotic systems use structured representations to describe objects, properties, relations, actions, and task constraints. Logic-based knowledge systems and robot knowledge bases have been used to store observed facts, infer missing information, and support semantic reasoning in robot tasks [53, 33]. For example, robots can infer likely object locations from qualitative spatial relations, container-location relations, or place-object co-occurrence knowledge [25, 15]. Commonsense and relational reasoning further connect symbolic knowledge with perceptual evidence in partially observable environments [37, 58, 7]. Ontologies and type hierarchies support consistent integration of new observations and generalization across related objects, actions, and task contexts [36, 40, 35].

Classical symbolic planning provides a computational layer on top of such knowledge representations. STRIPS and PDDL represent states as sets of symbolic facts and actions as preconditions and effects [12, 1]. Related planning representations such as FDR and domain-independent planners expose compact state variables and enable efficient search over symbolic transition systems [17, 18, 16, 5]. This structure is useful in robotics because high-level actions can be checked against symbolic preconditions and then grounded into executable robot skills. However, classical symbolic planning typically assumes that the current state is known, whereas real robot knowledge is often produced by noisy perception and partial observations.

Recent work has expanded robot knowledge beyond hand-coded symbolic rules using large language models and vision-language models. LLM-based planning systems use procedural and commonsense knowledge to propose task decompositions or action sequences, which are then

grounded with robot affordances, programmatic interfaces, or execution feedback [2, 20, 50, 30, 9]. Other systems construct richer environment knowledge from scene graphs and open-vocabulary spatial representations, allowing robots to reason about objects, places, and relations beyond a closed symbolic vocabulary [14, 28]. Retrieval-augmented methods use demonstrations, procedural knowledge, or task-relevant constraints to guide long-horizon robot planning [51, 60, 57, 54, 24]. Related approaches combine LLM-generated plans with symbolic constraints, spatial reasoning, or verification mechanisms to improve executability and consistency [19, 6, 32, 42, 43].

These systems show that modern knowledge sources can greatly broaden what robots know and how they generate plans. Nevertheless, both traditional knowledge bases and foundation-model-based systems often commit inferred knowledge as a proposal, prior, or verified fact rather than maintaining an explicit execution-time belief over alternative symbolic worlds. Their outputs may remain uncertain because of perception errors, incomplete grounding, missing observations, or hallucinated semantic assumptions. Our work keeps the interpretability and action structure of symbolic planning while representing uncertain local facts as a belief over possible symbolic states. This lets the planner reason about what it does not know before committing perceptual outputs or inferred facts to the knowledge base.

2.2 Belief-Space Planning and Symbolic Uncertainty

POMDPs provide a standard framework for decision making under partial observability [21]. Instead of assuming access to the true state, a POMDP maintains a belief distribution over possible states and selects actions to maximize expected return. Online methods such as POMCP scale this idea by using Monte Carlo tree search and particle beliefs over action-observation histories [49]. Related online POMDP methods have been applied to robot navigation, object search, manipulation, exploration, and interaction under uncertainty [52, 27, 56, 41].

These methods show that particle beliefs are effective when the robot must preserve multiple hypotheses about the world. However, many POMDP formulations operate over task-specific state variables rather than explicit symbolic knowledge bases. Conversely, symbolic planners expose action preconditions and logical structure, but often lack a principled belief representation. Prior work on symbolic planning under uncertainty addresses this gap by combining logical structure with probabilistic inference or POMDP-style planning [46, 59, 48]. Such methods demonstrate that symbolic structure can help organize uncertainty, but uncertainty is usually resolved through autonomous sensing, transition models, or planning abstraction.

Our setting adds a different source of information: human semantic feedback. Some ambiguities cannot be reliably resolved by additional robot sensing alone, such as visually ambiguous object states or uncertain task-relevant categories. We therefore maintain a symbolic particle belief over reachable frontier states and allow the robot to query a human when the belief is insufficiently confident. The answer directly filters the planner’s belief and updates the knowledge base.

2.3 Human Involvement in Robot Autonomy

Human involvement has long been used to improve robot reliability under uncertainty. In teleoperation and supervisory control, humans monitor the robot and intervene when needed [13, 10]. Human-on-the-loop systems preserve robot autonomy during nominal operation while relying on an operator to detect and correct failures [47, 39, 23, 34]. These systems can improve safety and

robustness, but they require sustained human attention. The operator must notice that a problem exists, interpret the situation, and decide how to intervene, which can increase workload and reduce scalability in long-horizon or repetitive tasks [11, 55].

Another approach is failure-triggered assistance, where the robot asks for help only after it fails, reaches a dead end, or detects an invalid state. This reduces continuous monitoring, but it treats human input as recovery rather than prevention. In tasks with irreversible or costly actions, asking after failure may be too late. A third approach is frequent or fixed querying, where the robot asks whenever a predefined uncertainty condition is met. This can improve task success, but it often produces unnecessary interactions because it does not distinguish uncertainty that matters for the current plan from uncertainty that is irrelevant.

Our work targets a proactive middle ground. The human is not asked to continuously supervise the robot, and the robot does not wait until failure. Instead, the robot estimates whether its current symbolic belief is confident enough to support planning. If not, it asks a targeted question about the symbolic fact expected to reduce the most uncertainty.

2.4 Uncertainty-Aware Querying

Recent work has studied robots and embodied agents that ask for help based on uncertainty. Active learning and active preference learning query humans to improve models or infer preferences [45]. Proactive robot behavior has also been shown to improve collaboration when the robot can anticipate when assistance is useful [3]. More recent uncertainty-aware systems use model confidence, language-model uncertainty, or conformal prediction to decide when to request human input [44, 31, 38]. These methods are important because they move beyond continuous supervision and show that selective querying can reduce unnecessary intervention.

However, many query systems focus on action-level uncertainty or model-output confidence. In such systems, the human may be asked which action the robot should take, whether a proposed action is acceptable, or whether a model prediction should be trusted. This differs from asking about the latent symbolic state that causes planning uncertainty. If the robot asks for an action, the human effectively participates in planning. If the robot asks for a state fact, the human provides information while the robot remains responsible for planning.

Our framework follows the second approach. The robot asks predicate-level questions about task-relevant symbolic facts, such as object properties, perceived categories, or local symbolic conditions. Human responses eliminate inconsistent hypotheses from the belief and update the knowledge base, after which planning proceeds autonomously. By integrating selective querying with symbolic belief maintenance and online planning, our approach treats human feedback as an information source for reducing planning uncertainty rather than as a mechanism for action selection.

3 Preliminaries

3.1 Symbolic Planning

Let $\mathcal{F}$ be a finite set of symbolic atoms. A symbolic state is a subset $s \subseteq \mathcal{F}$. An action $a \in \mathcal{A}$ is defined by preconditions, add effects, and delete effects:

$$a = \langle \text{pre}(a), \text{add}(a), \text{del}(a) \rangle.$$

The action is applicable in state s if $\text{pre}(a) \subseteq s$. Applying a produces

$$s' = (s \setminus \text{del}(a)) \cup \text{add}(a).$$

In robot domains, symbolic actions are grounded to executable skills such as detecting an object, navigating to a location, picking an object, placing an object, scanning a tomato, or discarding an item.

3.2 POMDPs and POMCP

A POMDP is a tuple $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, T, R, \Omega, O, \gamma \rangle$, where $\mathcal{S}$ is the state space, $\mathcal{A}$ the action space, $T(s'|s, a)$ the transition model, $R(s, a)$ the reward function, Ω the observation space, $O(o|s', a)$ the observation model, and γ the discount factor. Since the state is not directly observed, the robot maintains a belief $b(s)$ over states. After action a and observation o , the Bayesian belief update is

$$b'(s') = \eta O(o|s', a) \sum_{s \in \mathcal{S}} T(s'|s, a) b(s), \quad (1)$$

where η is a normalizer.

POMCP approximates online POMDP planning through Monte Carlo tree search over histories [49]. Each history node stores sampled particles and action values. During simulation, actions are selected using an upper-confidence rule,

$$a^* = \arg \max_a \left[V(ha) + c \sqrt{\frac{\log N(h)}{N(ha)}} \right], \quad (2)$$

where $V(ha)$ is the estimated value of taking action a after history h , $N(h)$ is the visit count of the history, and $N(ha)$ is the visit count of the action branch.

3.3 Human Feedback as Information

We treat a human response as an observation over a symbolic hypothesis. A hypothesis is a literal $q \in \mathcal{F}$ whose truth value is uncertain in the current belief. A query asks the human to verify q , producing an answer $y \in \{\text{true}, \text{false}\}$. The answer filters the belief by retaining states consistent with y . Thus, a query affects future planning by changing the posterior over symbolic worlds.

4 Method

4.1 Problem Formulation

The robot is given a symbolic domain, a set of executable skills, an initial knowledge state, and a task goal. Some facts are known with certainty, while local task facts must be inferred from observations. The robot selects task actions $a \in \mathcal{A}$ using POMCP. After each action, the observation may support several possible successor states rather than a single deterministic outcome. The system then computes belief uncertainty and requests human feedback only if the updated belief is not confident enough to commit a symbolic state to the knowledge base. The objective is to maximize task reward while reducing unnecessary feedback requests:

$$\max_{\pi} \mathbb{E}_{\pi} \left[\sum_{t=0}^T \gamma^t (R(s_t, a_t) - \lambda \mathbb{I}[\text{query}_t]) \right], \quad (3)$$

where query_t indicates that feedback is requested after the belief update at time t , and λ is the cost of human interaction.

4.2 Knowledge Base

The knowledge base is partitioned into domain information, static information, and local information:

$$\mathcal{KB} = \mathcal{I}^D \sqcup \mathcal{I}^S \sqcup \mathcal{I}^L. \quad (4)$$

Domain information encodes predicates, action schemas, task rules, and robot skills. Static information contains facts and numeric fluents that remain fixed during an episode, such as object identifiers, bin locations, or stem locations. Local information contains facts inferred during execution, such as object category, tomato state, object pose, and visibility. Uncertainty is concentrated in $\mathcal{I}^L$ because these facts depend on perception and partial observations.

4.3 Symbolic Frontier Belief

Let s^K denote the current committed knowledge state. Instead of constructing a global belief over all symbolic states, we construct a local belief over states reachable from s^K under the selected task action. For action a , the symbolic frontier is

$$\text{Frontier}(s^K, a) = \{s_1^*, \dots, s_j^*\}, \quad (5)$$

where each s_k^* is a reachable successor hypothesis consistent with the action model and the current knowledge base. The belief over the frontier is represented by weighted particles:

$$\mathcal{B} = \{(s_k^*, w_k)\}_{k=1}^j, \quad \sum_{k=1}^j w_k = 1. \quad (6)$$

After executing a and receiving observation o , each particle weight is updated by the observation model:

$$w'_k = w_k O(o | s_k^*, a), \quad \tilde{w}_k = \frac{w'_k}{\sum_{\ell=1}^j w'_\ell}. \quad (7)$$

This local frontier construction has two roles. First, it avoids enumerating irrelevant symbolic worlds. Second, it ties uncertainty to the action currently affecting task progress, which makes later human queries task-relevant.

4.4 Confidence and Query Trigger

We quantify uncertainty using the entropy of the normalized frontier belief:

$$H(\mathcal{B}) = - \sum_{k=1}^j \tilde{w}_k \log_2 \tilde{w}_k. \quad (8)$$

The confidence score is

$$\text{Conf}(\mathcal{B}) = \begin{cases} 1, & j = 1, \\ 1 - \frac{H(\mathcal{B})}{\log_2 j}, & j > 1. \end{cases} \quad (9)$$

If $\text{Conf}(\mathcal{B}) \geq \tau$, the robot commits the maximum-weight state to the knowledge base:

$$s^K \leftarrow \arg \max_{s_k^* \in \mathcal{B}} \tilde{w}_k. \quad (10)$$

If $\text{Conf}(\mathcal{B}) < \tau$, the robot requests human feedback. The threshold τ is therefore an autonomy-interaction parameter: lower values allow the robot to act with less certainty, while higher values increase interaction to obtain more reliable knowledge.

4.5 Query Selection

The system constructs candidate hypotheses from literals whose truth values differ across the frontier states. Let $\mathcal{C}$ be this set. For each candidate $c \in \mathcal{C}$, the belief is partitioned into states where c is true and false:

$$\mathcal{B}_c^+ = \{(s_k^*, \tilde{w}_k) \mid c \in s_k^*\}, \quad \mathcal{B}_c^- = \{(s_k^*, \tilde{w}_k) \mid c \notin s_k^*\}.$$

The selected query minimizes expected posterior entropy:

$$c^* = \arg \min_{c \in \mathcal{C}} [P(c)H(\mathcal{B}_c^+) + P(\neg c)H(\mathcal{B}_c^-)], \quad (11)$$

where $P(c) = \sum_{s_k^*: c \in s_k^*} \tilde{w}_k$ and $P(\neg c) = 1 - P(c)$. After the human answers, the belief is filtered:

$$\mathcal{B} \leftarrow \begin{cases} \mathcal{B}_{c^*}^+, & y = \text{true}, \\ \mathcal{B}_{c^*}^-, & y = \text{false}. \end{cases} \quad (12)$$

The remaining weights are renormalized. Querying repeats until the confidence exceeds τ or no informative hypothesis remains.

Figure Placeholder

fig:architecture

Figure 2: Detailed system architecture. The final figure should show the knowledge base, sensing and perception module, POMCP planning manager, feedback manager, user interface, and robot execution layer.

4.6 Planning and Execution Pipeline

At each step, the planning manager runs POMCP using the current belief and the symbolic action constraints. The selected task action is executed by a corresponding robot skill. Perception produces observations used to update the frontier belief. If the updated confidence is low, the feedback manager converts the selected symbolic hypothesis into a human-readable query, highlights the relevant object in the interface, receives a binary response, and returns the response to the planner. The knowledge base is updated only after the belief is sufficiently confident.

5 Experiments

The experiments are designed to answer four questions. First, does uncertainty-aware querying maintain task performance while reducing human interaction? Second, is the timing and content of queries important compared with random or fixed query policies? Third, how does the confidence threshold τ control the trade-off between autonomy and interaction? Fourth, does selective querying reduce human workload compared with continuous supervision or action-level intervention?

5.1 Domains

Waste sorting. The simulated waste sorting domain consists of a UR5 robot sorting objects into category-specific bins. Each object belongs to one of four categories: general waste, paper, can, or plastic. Categories are not fully known in advance and may be ambiguous due to occlusion or perception uncertainty. The robot must detect, pick, and place objects into the correct bins.

Tomato harvesting. The tomato harvesting domain contains tomatoes attached to stems. Each tomato may be ripe, unripe, or rotten. Ripe tomatoes should be harvested, rotten tomatoes should be discarded, and unripe tomatoes should be left untouched. The real-robot platform consists of a mobile base, a UR5e arm, an RGB-D camera, and a LiDAR sensor. This domain evaluates

Input: initial knowledge state s_0^K , threshold τ
Output: executed action sequence
 $s^K \leftarrow s_0^K$;
for $t = 0$ **to** T **do**
 $a_t \leftarrow \text{POMCP}(s^K, \mathcal{B})$;
 if $a_t = \emptyset$ **then**
 | **break**
 end
 execute a_t and receive observation o_{t+1} ;
 construct $\text{Frontier}(s^K, a_t)$ and update $\mathcal{B}$ using Eq. 7;
 compute $\text{Conf}(\mathcal{B})$ using Eq. 9;
 while $\text{Conf}(\mathcal{B}) < \tau$ **do**
 | select c^* using Eq. 11;
 | query human for c^* and receive y ;
 | filter $\mathcal{B}$ using Eq. 12;
 | recompute $\text{Conf}(\mathcal{B})$;
 end
 commit the most likely state to s^K using Eq. 10;
end

Algorithm 1: Uncertainty-Aware Planning with Human Querying

whether the framework can connect symbolic uncertainty, real perception, human feedback, and robot execution.

5.2 Baselines and Ablations

We compare the proposed method against the following strategies:

- **No Query:** the robot never asks for human feedback.
- **Always Query:** the robot queries at every uncertain decision point, approximating full human supervision.
- **Random Query:** the robot queries randomly with the same average query budget as the proposed method.
- **Failure-Triggered Query:** the robot asks only after reaching a failure, dead end, or invalid action.
- **Ours without Frontier:** the robot uses confidence thresholding but selects queries without action-relevant frontier structure.
- **Ours:** the full uncertainty-aware frontier query framework.




Figure 3: Experimental domains. The final figure should show the simulated waste sorting setup and the real tomato harvesting robot setup.

Table 1: System evaluation template. Replace placeholders after final experiments.

Domain	Method	Success (%)	Reward	Queries	Time (s)
Waste	No Query	—	—	—	—
	Always Query	—	—	—	—
	Random Query	—	—	—	—
	Failure-Triggered	—	—	—	—
	Ours w/o Frontier	—	—	—	—
	Ours	—	—	—	—
Tomato	No Query	—	—	—	—
	Always Query	—	—	—	—
	Random Query	—	—	—	—
	Failure-Triggered	—	—	—	—
	Ours w/o Frontier	—	—	—	—
	Ours	—	—	—	—

5.3 Metrics

We report task success rate, cumulative reward, average number of queries, query ratio, task completion time, and failure type. For user studies, we additionally report subjective workload using NASA-TLX and interaction preference ratings. For real-robot evaluation, we report scenario-level success, number of queries, intervention timing, and representative failure cases.

5.4 System Evaluation

The expected result is that Always Query achieves high success but incurs the largest interaction cost, No Query reduces interaction but fails under unresolved ambiguity, and Random Query underperforms because it often spends queries on non-critical states. The full method should approach the success of Always Query with far fewer queries by targeting uncertainty that is relevant to planning.

Figure Placeholder

fig:tau

Figure 4: Threshold sensitivity for τ . The final figure should include aggregate and per-scenario plots for waste sorting and tomato harvesting.

5.5 Threshold Analysis

The confidence threshold τ controls the autonomy-interaction trade-off. Low τ values make the robot more willing to commit uncertain beliefs, reducing queries but increasing failure risk. High τ values make the robot ask more often, improving confidence but increasing workload. The intended analysis should show success rate, reward, and query count across τ values in both domains. This analysis should be moved into the main paper rather than treated only as appendix material, because τ is the key parameter governing selective autonomy.

5.6 Real-Robot Tomato Evaluation

The real-robot study should include repeated trials across multiple tomato scenarios. Each scenario should specify the initial arrangement, tomato states, visual ambiguities, and expected task outcome. The evaluation should report whether each query prevented a wrong harvest, wrong discard, or missed target. A qualitative case study should compare the same scene under No Query and Ours, showing how a human answer changes the belief and the subsequent plan.

5.7 User Study

The user study evaluates whether system-initiated state queries reduce workload compared with continuous supervision and action-level assistance. Participants observe or interact with the robot through the interface while the robot executes task scenarios containing perceptual ambiguity. We compare at least three interaction conditions: **Always Query**, where the user is repeatedly asked to verify uncertain information; **Failure-Triggered Query**, where the user is asked only after the robot reaches an invalid or failed state; and **Ours**, where the system asks targeted state-level questions when entropy-based confidence is low.

We measure subjective workload using NASA-TLX, along with the number of user responses, response time, perceived interruption, trust, and preference. The main hypothesis is that the proposed method reduces workload because the user is not required to continuously monitor the robot or make action-level planning decisions. Instead, the user answers targeted questions about symbolic state facts, while the robot remains responsible for belief update and replanning.

Table 2: User study template. Replace placeholders after collecting user-study results.

Condition	NASA-TLX	Responses	Response Time	Preference
Always Query	—	—	—	—
Failure-Triggered	—	—	—	—
Action-Level Query	—	—	—	—
Ours	—	—	—	—

6 Discussion

The central claim of this paper is that asking a human should be treated as part of robot planning, not as an external recovery behavior. This matters because the usefulness of a query depends on the robot’s current belief and on the task consequences of remaining uncertain. A query about a visually ambiguous tomato is valuable if the answer changes whether the robot harvests, discards, or ignores it. The same query may be unnecessary if the tomato is irrelevant to the current goal or if another action can resolve the uncertainty autonomously.

The frontier belief representation is intended to make this connection explicit. By constructing uncertainty around states reachable from the current knowledge state and selected action, the robot avoids asking about facts that do not affect near-term planning. The entropy-based confidence score determines when the robot lacks enough information to commit a state, and the expected-entropy query rule determines which symbolic fact would most reduce the ambiguity.

Several limitations should be discussed clearly. First, the framework assumes that human answers are reliable or at least more reliable than the robot’s uncertain observation. Future work should model noisy human feedback. Second, the current query form is binary and predicate-level. More expressive questions may reduce the number of interactions but increase user burden. Third, the planner uses a local frontier belief rather than a full global posterior. This improves tractability but may miss dependencies outside the immediate action-relevant frontier. Finally, the real-robot evaluation must show repeated trials rather than isolated demonstrations for the method to be convincing as a robotics contribution.

7 Conclusion

We presented a knowledge-based belief-space planning framework that integrates proactive human querying into robot decision making. The robot maintains uncertainty over symbolic possible worlds, uses POMCP for online planning, and requests human feedback only when the observation-updated belief remains insufficiently confident. By selecting queries that are expected to reduce entropy over task-relevant frontier states, the system decides when and what to ask after ambiguous observations. This formulation uses human feedback as a belief-refinement mechanism rather than as action-level control, enabling robots to reduce unnecessary supervision while resolving ambiguity before it causes task failure. The user study further evaluates whether this selective state-level interaction reduces subjective workload compared with continuous supervision or action-level assistance.

References

- [1] Constructions Aeronautiques, Adele Howe, Craig Knoblock, ISI Drew McDermott, Ashwin Ram, Manuela Veloso, Daniel Weld, David Wilkins Sri, Anthony Barrett, Dave Christianson, et al. PDDL—the planning domain definition language. *Technical Report, Tech. Rep.*, 1998.
- [2] Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, et al. Do as i can, not as i say: Grounding language in robotic affordances, 2022.
- [3] Jimmy Baraglia, Maya Cakmak, Yukie Nagai, Rajesh PN Rao, and Minoru Asada. Efficient human-robot collaboration: when should a robot take initiative? *The International Journal of Robotics Research*, 36(5-7):563–579, 2017.
- [4] Jessie YC Chen and Michael J Barnes. Human–agent teaming for multirobot control: A review of human factors issues. *IEEE Transactions on Human-Machine Systems*, 44(1):13–29, 2014.
- [5] Augusto B Corrêa and Giuseppe De Giacomo. Lifted planning: Recent advances in planning using first-order representations. In *IJCAI*, pages 8010–8019, 2024.
- [6] Zhirui Dai, Arash Asgharivaskasi, Thai Duong, Shusen Lin, Maria-Elizabeth Tzes, George Pappas, and Nikolay Atanasov. Optimal scene graph planning with large language model guidance. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 14062–14069, 2024.
- [7] Angel Daruna, Weiyu Liu, Zsolt Kira, and Sonia Chetnova. RoboCSE: Robot common sense embedding. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 9777–9783, 2019.
- [8] Randall Davis, Howard Shrobe, and Peter Szolovits. What is a knowledge representation? *AI magazine*, 14(1):17–17, 1993.
- [9] Yan Ding, Xiaohan Zhang, Chris Paxton, and Shiqi Zhang. Task and motion planning with large language models for object rearrangement. In *Proc. of IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 2086–2092, 2023.
- [10] Anca D Dragan, Kenton CT Lee, and Siddhartha S Srinivasa. Legibility and predictability of robot motion. In *International Conference on Human-Robot Interaction (HRI)*, pages 301–308. IEEE, 2013.
- [11] Mica R Endsley. Automation and situation awareness. In *Automation and human performance*, pages 163–181. CRC Press, 2018.
- [12] Richard E Fikes and Nils J Nilsson. Strips: A new approach to the application of theorem proving to problem solving. *Artificial intelligence*, 2(3-4):189–208, 1971.
- [13] Michael A Goodrich and Alan C Schultz. Human–robot interaction: a survey. *Foundations and trends® in human–computer interaction*, 1(3):203–275, 2008.

- [14] Qiao Gu, Ali Kuwajerwala, Sacha Morin, Krishna Murthy Jatavallabhula, Bipasha Sen, Aditya Agarwal, Corban Rivera, William Paul, Kirsty Ellis, Rama Chellappa, et al. Conceptgraphs: Open-vocabulary 3d scene graphs for perception and planning. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 5021–5028. IEEE, 2024.
- [15] Shoichi Hasegawa, Akira Taniguchi, Yoshinobu Hagiwara, Lotfi El Hafi, and Tadahiro Taniguchi. Inferring place-object relationships by integrating probabilistic logic and multimodal spatial concepts. In *Proc. of IEEE/SICE International Symposium on System Integration (SII)*, pages 1–8, 2023.
- [16] Malte Helmert. The fast downward planning system. *Journal of Artificial Intelligence Research*, 26:191–246, 2006.
- [17] Malte Helmert. Concise finite-domain representations for pddl planning tasks. *Artificial Intelligence*, 173(5-6):503–535, 2009.
- [18] Jörg Hoffmann. Ff: The fast-forward planning system. *AI magazine*, 22(3):57–57, 2001.
- [19] Wenlong Huang, Chen Wang, Ruohan Zhang, Yunzhu Li, Jiajun Wu, and Li Fei-Fei. Vox-Poser: Composable 3d value maps for robotic manipulation with language models. In *Proc. of Conference on Robot Learning (CoRL)*, pages 540–562, 2023.
- [20] Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, et al. Inner monologue: Embodied reasoning through planning with language models, 2022.
- [21] Leslie Pack Kaelbling, Michael L Littman, and Anthony R Cassandra. Planning and acting in partially observable stochastic domains. *Artificial intelligence*, 101(1-2):99–134, 1998.
- [22] Krishnanand N Kaipa, Akshaya S Kankanhalli-Nagendra, Nithyananda B Kumbla, Shaurya Shriyam, Srudeep Somnaath Thevendria-Karthic, Jeremy A Marvel, and Satyandra K Gupta. Addressing perception uncertainty induced failure modes in robotic bin-picking. *Robotics and Computer-Integrated Manufacturing*, 42:17–38, 2016.
- [23] Michael Kelly, Chelsea Sidrane, Katherine Driggs-Campbell, and Mykel J Kochenderfer. Hg-dagger: Interactive imitation learning with human experts. In *2019 International Conference on Robotics and Automation (ICRA)*, pages 8077–8083. IEEE, 2019.
- [24] Nishanth Kumar, William Shen, Fabio Ramos, Dieter Fox, Tomás Lozano-Pérez, Leslie Pack Kaelbling, and Caelan Reed Garrett. Open-world task and motion planning via vision-language model generated constraints. *IEEE Robotics and Automation Letters*, 2026.
- [25] Lars Kunze, Keerthi Kumar Doreswamy, and Nick Hawes. Using qualitative spatial relations for indirect object search. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 163–168, 2014.
- [26] Hanna Kurniawati and Vinay Yadav. An online pomdp solver for uncertainty planning in dynamic environment. In *Robotics Research: The 16th International Symposium ISRR*, pages 611–629. Springer, 2016.

- [27] Mikko Lauri and Risto Ritala. Planning for robotic exploration based on forward simulation. *Robotics and Autonomous Systems*, 83:15–31, 2016.
- [28] Alessandro Leanza, Angelo Moroncelli, Giuseppe Vizzari, Francesco Braghin, Loris Roveda, and Blerina Spahiu. Conceptbot: Enhancing robot’s autonomy through task decomposition with large language models and knowledge graph. *arXiv preprint arXiv:2509.00570*, 2025.
- [29] Jan Leusmann, Chao Wang, Michael Gienger, Albrecht Schmidt, and Sven Mayer. Understanding the uncertainty loop of human-robot interaction. *arXiv preprint arXiv:2303.07889*, 2023.
- [30] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 9493–9500, 2023.
- [31] Kaiqu Liang, Zixu Zhang, and Jaime F Fisac. Introspective planning: Aligning robots’ uncertainty with inherent task ambiguity. *Advances in Neural Information Processing Systems*, 37:71998–72031, 2024.
- [32] Yuchen Liu, Luigi Palmieri, Sebastian Koch, Ilche Georgievski, and Marco Aiello. DELTA: Decomposed efficient long-term robot task planning using large language models. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 10995–11001, 2025.
- [33] Zhiyu Liu, Meng Jiang, and Hai Lin. A graph-based spatial temporal logic for knowledge representation and automated reasoning in cognitive robots, 2020.
- [34] Ajay Mandlekar, Danfei Xu, Roberto Martín-Martín, Yuke Zhu, Li Fei-Fei, and Silvio Savarese. Human-in-the-loop imitation learning using remote teleoperation. *arXiv preprint arXiv:2012.06733*, 2020.
- [35] Sumaira Manzoor, Yuri Goncalves Rocha, Sung-Hyeon Joo, Sang-Hyeon Bae, Eun-Jin Kim, Kyeong-Jin Joo, and Tae-Yong Kuc. Ontology-based knowledge representation in robotic systems: A survey oriented toward applications. *Applied Sciences*, 11(10):4324, 2021.
- [36] Alex Mitrevsk, Paul G. Plöger, and Gerhard Lakemeyer. Ontology-assisted generalisation of robot action execution knowledge. In *Proc. of IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 6763–6770, 2021.
- [37] Tiago Mota, Mohan Sridharan, and Aleš Leonardis. Integrated commonsense reasoning and deep learning for transparent decision making in robotics. *SN Computer Science*, 2(4):242, 2021.
- [38] James F Mullen and Dinesh Manocha. Lbap: Improved uncertainty alignment of llm planners using bayesian inference. In *2025 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 18716–18723. IEEE, 2025.

- [39] Saeid Nahavandi. Trusted autonomy between humans and robots: Toward human-on-the-loop in robotics and autonomous systems. *IEEE Systems, Man, and Cybernetics Magazine*, 3(1):10–17, 2017.
- [40] Alberto Olivares-Alarcos, Daniel Beßler, Alaa Khamis, Paulo Goncalves, Maki K. Habib, Julita Bermejo-Alonso, Marcos Barreto, Mohammed Diab, Jan Rosell, João Quintas, et al. A review and comparison of ontology-based approaches to robot autonomy. *The Knowledge Engineering Review*, 34:e29, 2019.
- [41] Joni Pajarinen, Jens Lundell, and Ville Kyrki. POMDP planning under object composition uncertainty: Application to robotic manipulation. *IEEE Transactions on Robotics*, 39(1):41–56, 2023.
- [42] David Paulius, Alejandro Agostini, Benedict Quartey, and George Konidaris. Bootstrapping object-level planning with large language models. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 16233–16239, 2025.
- [43] Benedict Quartey, Eric Rosen, Stefanie Tellex, and George Konidaris. Verifiably following complex robot instructions with foundation models. In *Proc. of IEEE International Conference on Robotics and Automation (ICRA)*, pages 1–8, 2025.
- [44] Allen Z. Ren, Anushri Dixit, Alexandra Bodrova, Sumeet Singh, Stephen Tu, Noah Brown, Peng Xu, Leila Takayama, Fei Xia, Jake Varley, Zhenjia Xu, Dorsa Sadigh, Andy Zeng, and Anirudha Majumdar. Robots that ask for help: Uncertainty alignment for large language model planners. In *Proceedings of the Conference on Robot Learning (CoRL)*, 2023.
- [45] Dorsa Sadigh, Anca D. Dragan, S. Shankar Sastry, and Sanjit A. Seshia. Active preference-based learning of reward functions. In *Proceedings of Robotics: Science and Systems (RSS)*, 2017.
- [46] Scott Sanner and Kristian Kersting. Symbolic dynamic programming for first-order pomdps. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 24, pages 1140–1146, 2010.
- [47] Paul Scharre. Introduction to autonomy in weapon systems. 2015.
- [48] Sergio A Serrano, Elizabeth Santiago, Jose Martinez-Carranza, Eduardo F Morales, and L Enrique Sucar. Knowledge-based hierarchical pomdps for task planning. *Journal of Intelligent & Robotic Systems*, 101(4):82, 2021.
- [49] David Silver and Joel Veness. Monte-carlo planning in large pomdps. *Advances in neural information processing systems*, 23, 2010.
- [50] Ishika Singh, Valts Blukis, Arsalan Mousavian, Ankit Goyal, Danfei Xu, Jonathan Tremblay, Dieter Fox, Jesse Thomason, and Animesh Garg. ProgPrompt: Program generation for situated robot task planning using large language models. *Autonomous Robots*, 47(8):999–1012, 2023.

- [51] Shivam Singh, Karthik Swaminathan, Nabanita Dash, Ramandeep Singh, Snehasis Banerjee, Mohan Sridharan, and Madhava Krishna. Adaptbot: Combining llm with knowledge graphs and human input for generic-to-specific task decomposition and knowledge refinement. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 4345–4351. IEEE, 2025.
- [52] Adhiraj Somani, Nan Ye, David Hsu, and Wee Sun Lee. Despot: Online pomdp planning with regularization. *Advances in neural information processing systems*, 26, 2013.
- [53] Moritz Tenorth and Michael Beetz. Knowledge processing for autonomous robot control. In *Proc. of AAAI Spring Symposium: Designing Intelligent Robots*, 2012.
- [54] Zheng Wang, Shu Xian Teo, Jun Jie Chew, and Wei Shi. InstructRAG: Leveraging retrieval-augmented generation on instruction graphs for llm-based task planning. In *Proc. of International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*, 2025.
- [55] Choon Yue Wong and Gerald Seet. Workload, awareness and automation in multiple-robot supervision. *International Journal of Advanced Robotic Systems*, 14(3):1729881417710463, 2017.
- [56] Yuchen Xiao, Sammie Katt, Andreas ten Pas, Shengjian Chen, and Christopher Amato. On-line planning for target object search in clutter under partial observability. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, pages 8241–8247. IEEE, 2019.
- [57] Weiye Xu, Min Wang, Wengang Zhou, and Houqiang Li. P-RAG: Progressive retrieval augmented generation for planning on embodied everyday task. In *Proc. of ACM International Conference on Multimedia (MM)*, pages 6969–6978, 2024.
- [58] Shiqi Zhang, Piyush Khandelwal, and Peter Stone. iCORPP: Interleaved commonsense reasoning and probabilistic planning on robots. *Robotics and Autonomous Systems*, 174:104613, 2024.
- [59] Shiqi Zhang and Peter Stone. Corpp: Commonsense reasoning and probabilistic planning, as applied to dialog with a mobile robot. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 29, 2015.
- [60] Jiacong Zhou, Jiaxu Miao, Xianyun Wang, and Jun Yu. Enhancing llm planning for robotics manipulation through hierarchical procedural knowledge graphs. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025.

A Appendix

A.1 Experiment Settings

This appendix should list the exact planner parameters, perception parameters, domain parameters, and scenario descriptions used in the waste sorting and tomato harvesting experiments.

A.2 Additional Threshold Results

Additional per-scenario threshold plots and raw statistics for τ should be placed here after the main threshold analysis is summarized in Sec. 5.